

What is the Internet?



Question

Here is a typical software engineering interview question:

"What happens when you type google.com into your browser's address box and press enter?"

Reference: <https://github.com/alex/what-happens-when>

Answer

A bird's-eye view gives insight into the "full stack"

- G key is pressed
- The "enter" key bottoms out
- Browser parses the URL
- Browser looks up the DNS
- Browser makes a connection
- Request and response
- Browser parses response
- Browser renders page

Learning objectives

- Understand and explain what happens when you type google.com into your browser's address box and press enter
- Gain an appreciation of complexity hidden beneath the hood of everyday actions
- Learn important terms and concepts in "full stack" web development, which we will come up again in more depth later

The "g" key is pressed

What you see:

- Autocomplete

What you don't see:

- Lookups for suggestions

The "enter" key bottoms out

What you see:

- Things start to happen. Your browser might blank out, or start refreshing
- Progress bars, spinners
- "X" instead of refresh symbol ↺

The "enter" key bottoms out (cont.)

What you don't see:

- State of the browser changes. Goes into this in-between loading state.
- The `WM_KEYDOWN` or `KeyDown` event
- The keycode `13`, which stands for the "enter" key, is transmitted to the computer (USB, PS/2 interface, or a "virtualized" keyboard, like from a touchscreen)

Browser parses the URL

- Browser looks for the **protocol** "http", the **hostname** "google.com" and the **resource** "/".
- Browser decides if this is a URL or a search term. Browsers will default to searching if what's entered in the address bar isn't a valid URL.

Browser parses the URL (cont.)

- Checks if website needs to be redirected via the HTTP Strict Transport Security (HSTS) list.
- Converts non-ASCII Unicode characters in hostname. Only certain characters are valid for hostnames, **a-z**, **A-Z**, **0-9**, **-**, and **.**

Browser looks up the DNS

- First, it checks the `cache`.
- If it can't find it in the cache, it'll `gethostbyname`, which first checks the `hosts` file (OS dependent).
- If it can't find the host in the `hosts` file, it will look for it in the DNS server, which is determined by the network (router, modem, etc.)

Browser makes a connection

- Now that it has the IP address, it will open a **socket**
- What's important here is the IP address and the **port**, for example a router is typically configured to **192.168.1.1:80**. The IP address is **192.168.1.1**, and **80** is the port number.
- Another example of an IP address is **127.0.0.1** which is typically your **localhost**.

Request

- If the protocol is secure (HTTPS), it will do a TLS (Transport Layer Security) handshake
- For both HTTP and HTTPS, it will do an HTTP **GET** request
- For each **request**, there is a **response**


Response

- A typical response is a **200 OK**, which means everything was good
- Other responses include **301 Moved Permanently**, which is a redirect, **404 Not Found**, and **500 Internal Server Error**.
- The first digit usually signifies what kind of code it is (e.g. **2xx**, **3xx**, **4xx**, or **5xx**).

Browser parses resources

- The response comes along with **resources**, which include all the data that a browser needs to render a page: HTML, CSS, JavaScript, images, etc.
- Browser has to parse each of these resources to construct a **tree** known as the **DOM** (document object model).

Browser parses resources (cont.)




Browser parses resources (cont.)

- The DOM represents each of the HTML **elements** on a page and their hierarchical relationships. Each element also comes with a set of **attributes**.
- An element is an individual component, whereas an attribute is a modifier specific to that individual component (e.g. `<p>`
`</p>` is an element and `<p`
`class="title"></p>` is a `p` element with a class attribute set to `title`)

Browser renders page

- It goes through each of the nodes in the DOM tree and calculates the sizes of them on the page
- These calculations are dynamically determined by a bunch of factors: CSS box model (margins, borders, padding), viewport, CSS position (float, absolute, relative), etc.
- Does actual graphic rendering of page (often involves GPU commands through Direct3D or OpenGL)

Summary



- A ton of things happen when you type google.com into your browser's address box and press enter

Summary (cont.)

- Many moving parts, interdependent pieces of a larger system
- Most of this is abstracted away in different layers, which allows you to think of only the problem at hand

Next: The layer we're going to start is at the **command line**. This is how developers manipulate the files that are needed to run do any of the things we just mentioned

Reference: <https://github.com/alex/what-happens-when>